




# Braver Cat


**Brain** and **Liver cancer** segmentation via **T**ransfer Learning

Daniele Solombrino, Emanuele Volanti

# Project idea



- Basic: **brain** tumor segmentation
  - **End-to-end** train
- Advanced: **liver** tumor segmentation
  - **Transfer Learning**  
using NN trained in brain step



# Motivations



- Work on a typical Computer Vision task
  - Segmentation
- Experience a full Deep Learning pipeline
  - Data gathering
  - Data pre-processing
  - Model engineering
  - Model evaluation
  - Result analysis
- Very small datasets for very specific pathologies

# Paper hunt




- Literature research process
  - Find interesting papers in **surveys**
  - **Deep dive** on specific papers
- Literature review results
  - “[Cheplygina et al., 2019](#)” → shallow models 👎
  - “[A Survey on DL in Medical Image Analysis](#)” → Deep Learning ❤️
  - “[GANs for Medical Image Analysis](#)” → computationally expensive 🤔
  - “[Adversarial Methods in Medical Analysis](#)” → computationally expensive 🤔
  - “[Brain Tumor Segmentation with DNNs](#)” → DL 👍, computationally feasible 👍, novelties 👍

# TwoPathCNN






- Multiple filter scales adoption → **different scale** representations
- Combining them → CNN learns much **more** feature **interactions**
- Medical domain: non-local, potentially unknown **tissue relationships**



# CNN concatenation




- TwoPathCNN: multi-scale interactions, single input...
- ... what about **multiple inputs**?
- Boosts TwoPathCNN benefits
- Multiple concatenation strategies
  - Task-dictated



# Dataset and patched inputs



- Dataset → brain scans w/ segmentation masks
- **4-channelled** inputs → 4 CT modalities
- 5 labels → possible tumor stages
- For every **pixel p**, model:
  - Sees **differently-scaled** patches centered on **p**
  - Predicts segmentation label for pixel **p**



# Dual-stage training as unbalanced data remedy



- **Unbalancement** towards “non-tumor” label 👎
- 💡 Idea: **dual-stage training**
- 1st stage → balanced data
- 2nd stage → unbalanced data
  - Resume from 1st stage w/ all but last layer frozen







# Convolutional output layer



- Typical CNNs use a **fully-connected** output layer
  - Needs 1D **flattening** → wastes time and memory 🙄
  - Can **not** be **parallelized** 🙄
- 💡 Idea: 1x1 **conv** layer, w/  $c$  channels
  - $c \rightarrow$  number of classes
  - Way easier to **parallelize** 👍
  - Weight sharing 👍

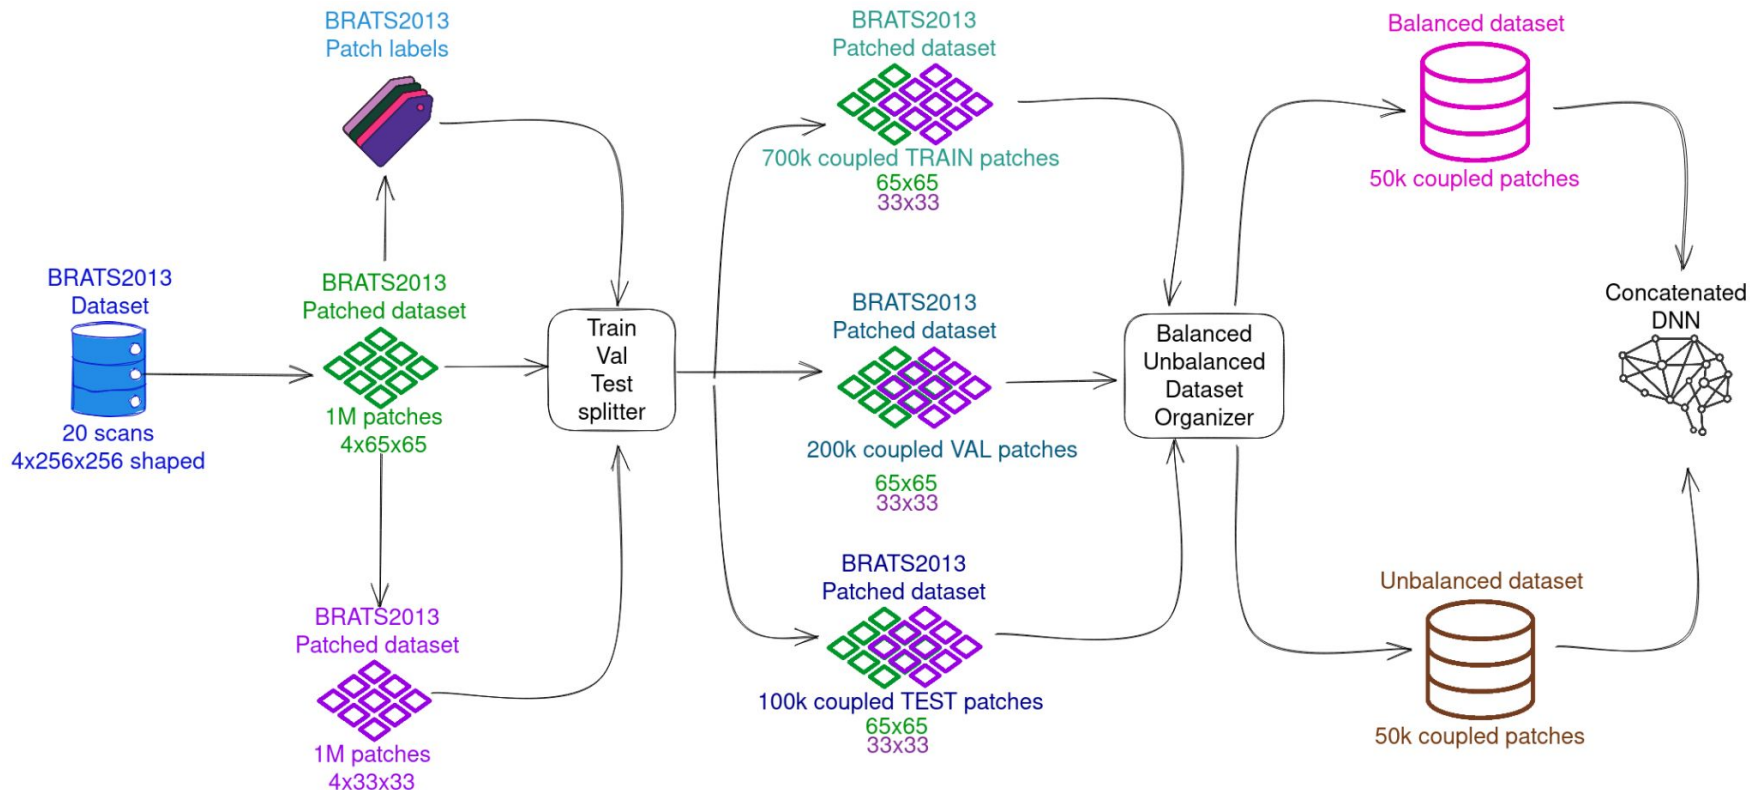



# Challenges




- Complex pre-processing
- Custom Neural Networks from scratch
- Hyperparameter tuning
- Transfer Learning
  - What layers to freeze?
  - Input channels mismatch

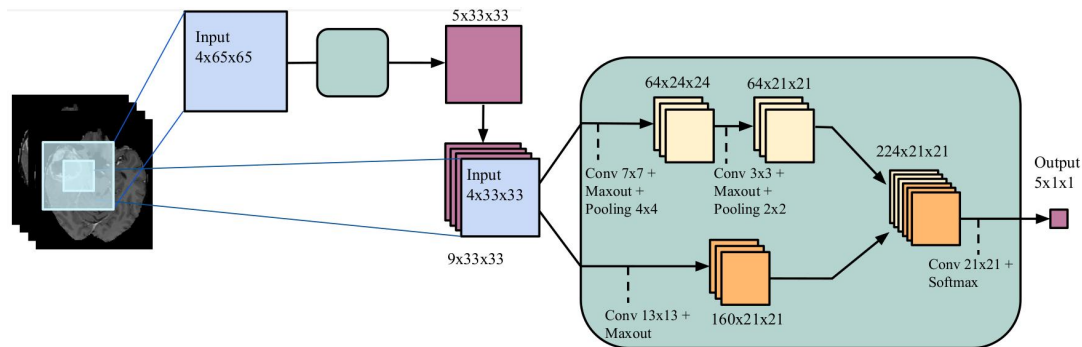
# Non-trivial pre-processing




# Custom Neural Networks from scratch




- Defined **PyTorch** modules for TwoPathCNN
- **Maxout** activation → **no** out of the box PyTorch **support**
- Implemented InputCascadeCNN model → two **concatenated** TwoPathCNN



# Hyperparameter tuning



- Paper lists hyperparameters...
- ... but
  - No values disclosed, or
  - Proposed values **perform badly**
- ... so we worked on
  - Weights init method, dropout
  - Optimizer, LR w/ decay, momentum
  - Regularization
  - Input normalization, batch size





# End-to-end dual-stage training





- 1st stage
  - ~2.2M trainable parameters
  - 600 epochs
  - ~5h, 30s per epoch
  - **>90%** train, val and test acc
- 2nd stage
  - ~600k trainable parameters
  - 300 epochs
  - 40 mins, 4s per epoch
  - **>92%** train, val and test acc

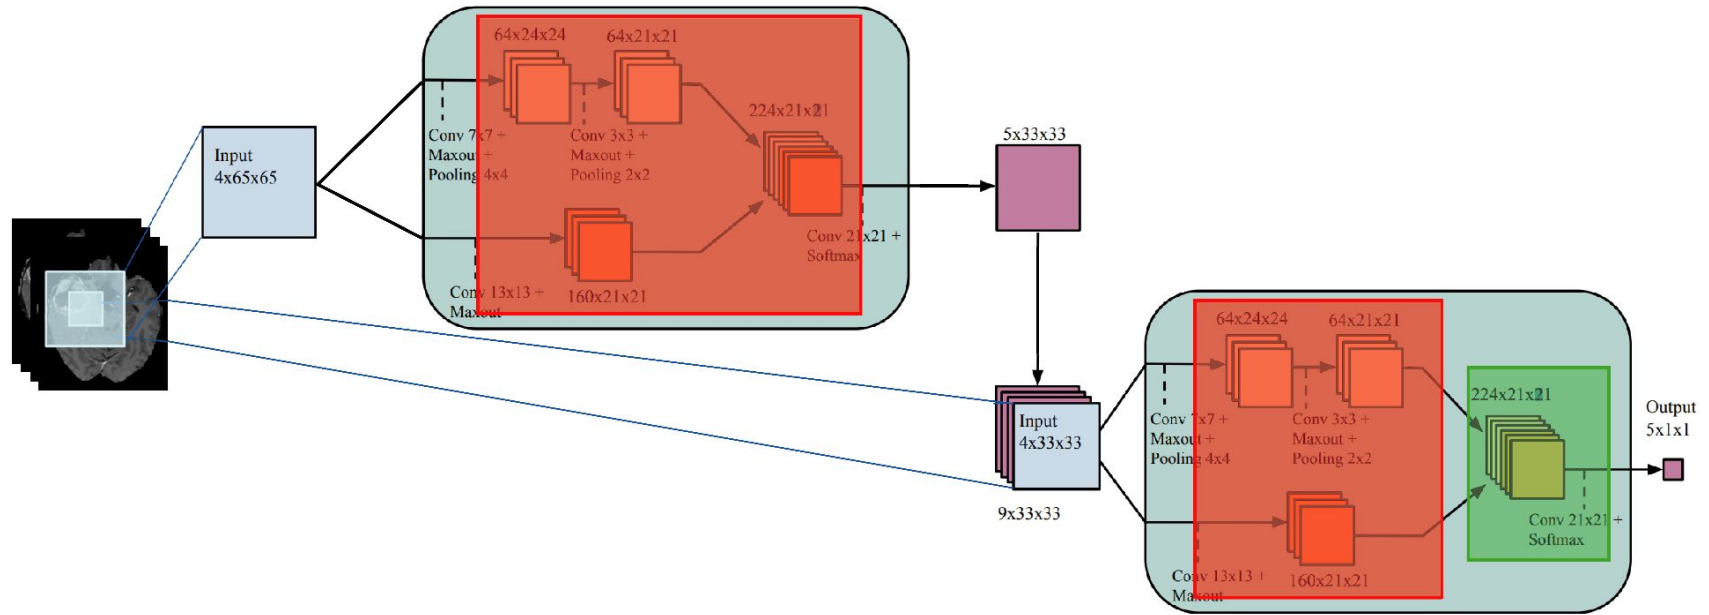

accuracy 2



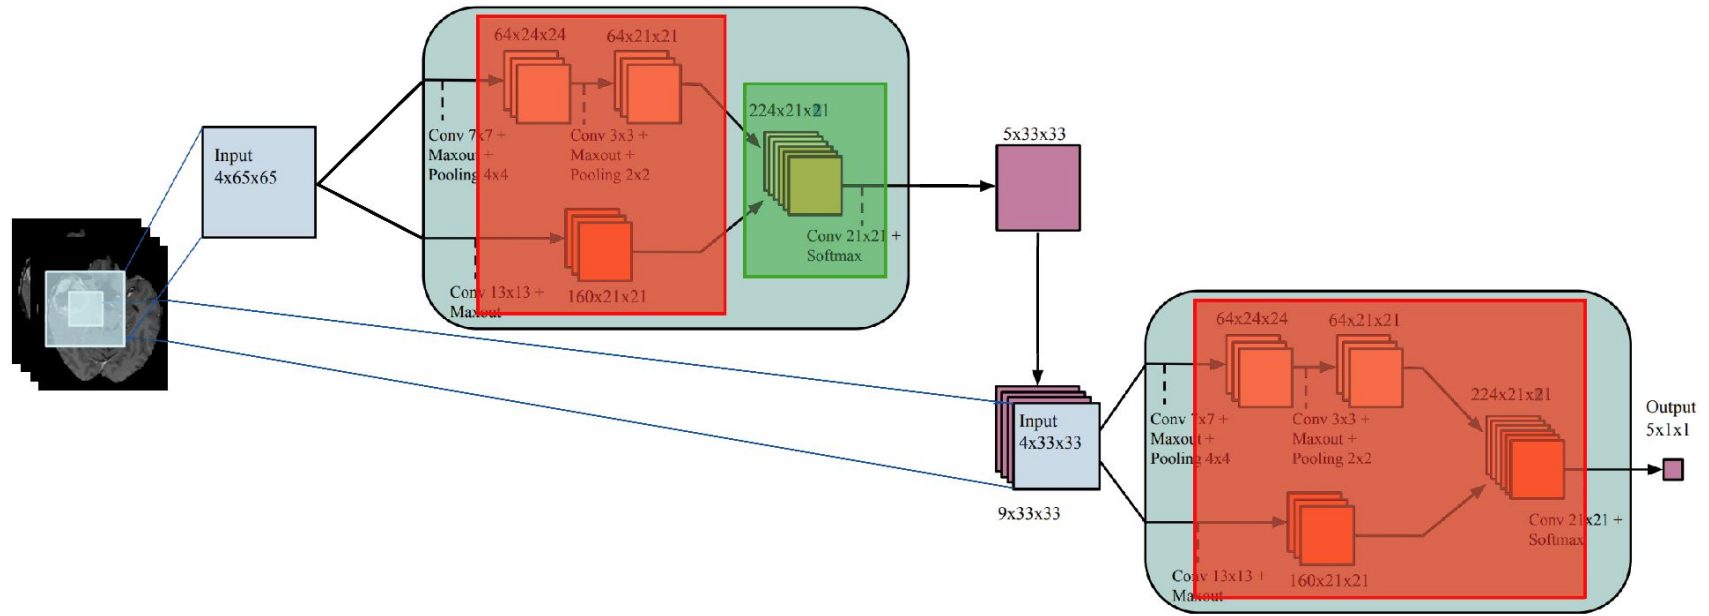

loss 2



# Transfer Learning: local scale

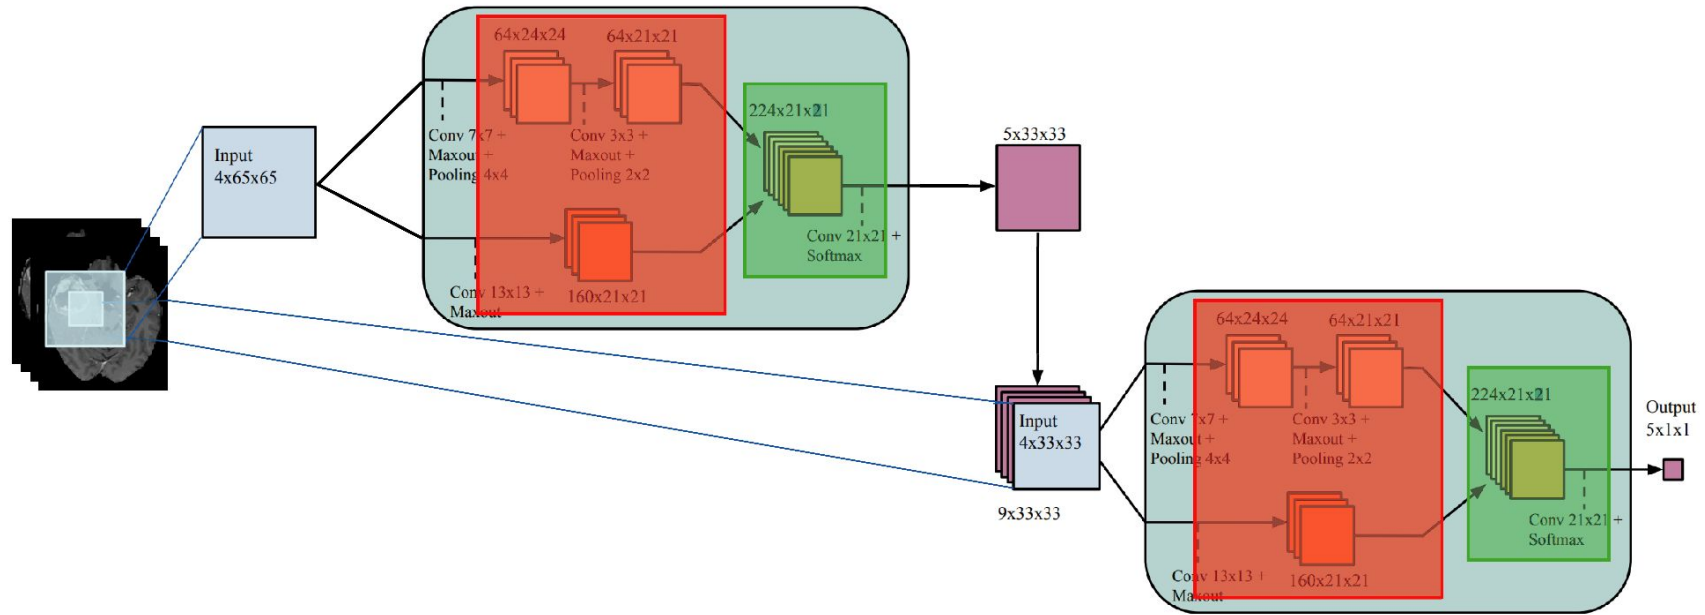



# Transfer Learning: global scale






# Transfer Learning: local and global scales



# Transfer Learning: results



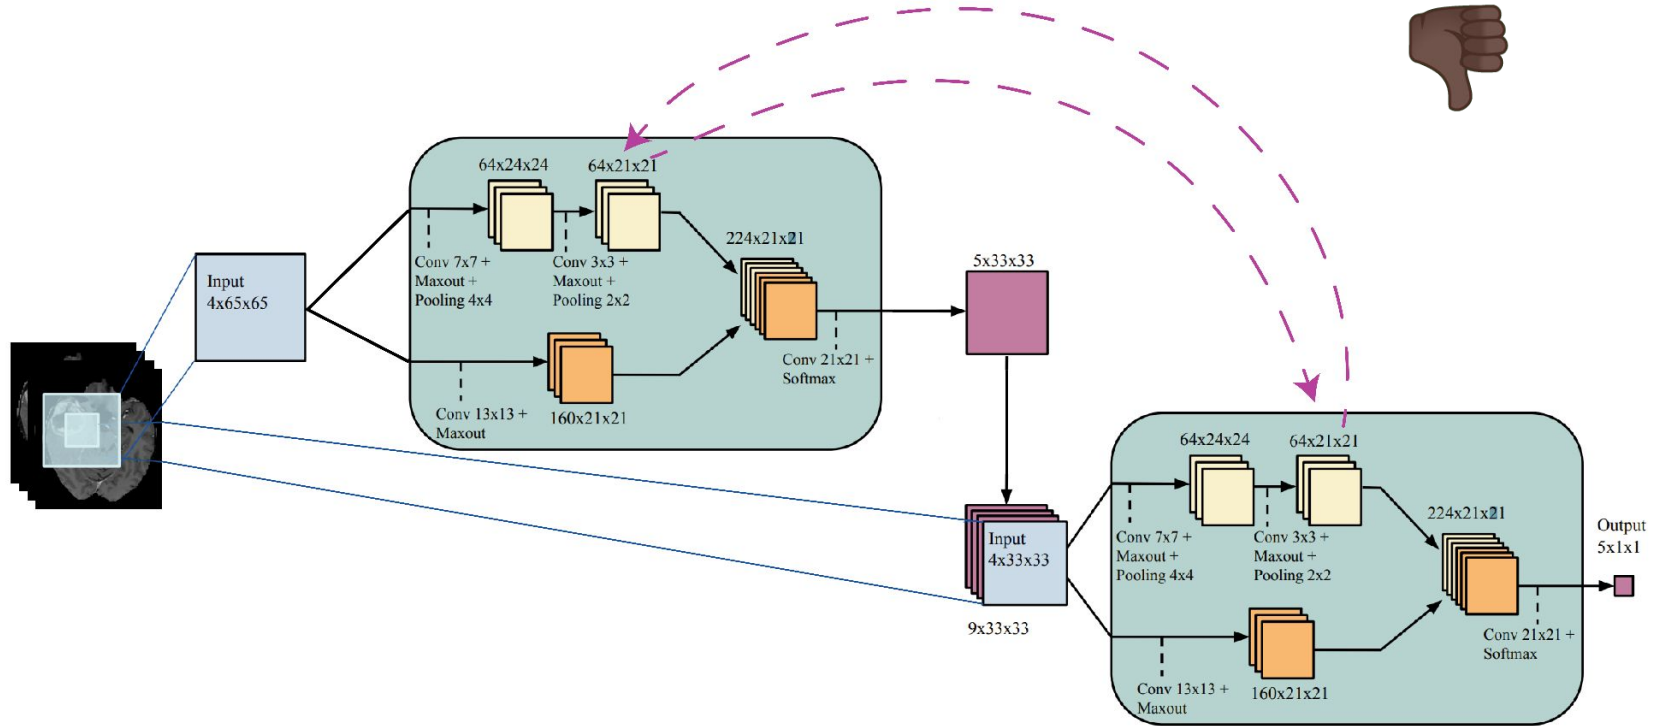
| Transfer Learning setup | # trainable params | Epoch train time | Train accuracy | Validation accuracy |
|-------------------------|--------------------|------------------|----------------|---------------------|
| Local scale             | 600k               | 4s               | 52%            | 48%                 |
| Global scale            | 600k               | 4s               | 66%            | 61%                 |
| Global & Local scale    | 1.2M               | 9.5s             | 87%            | 85%                 |
| E2E                     | 2.2M               | 30s              | 89%            | 86%                 |

# Hypothesis: parameter pruning

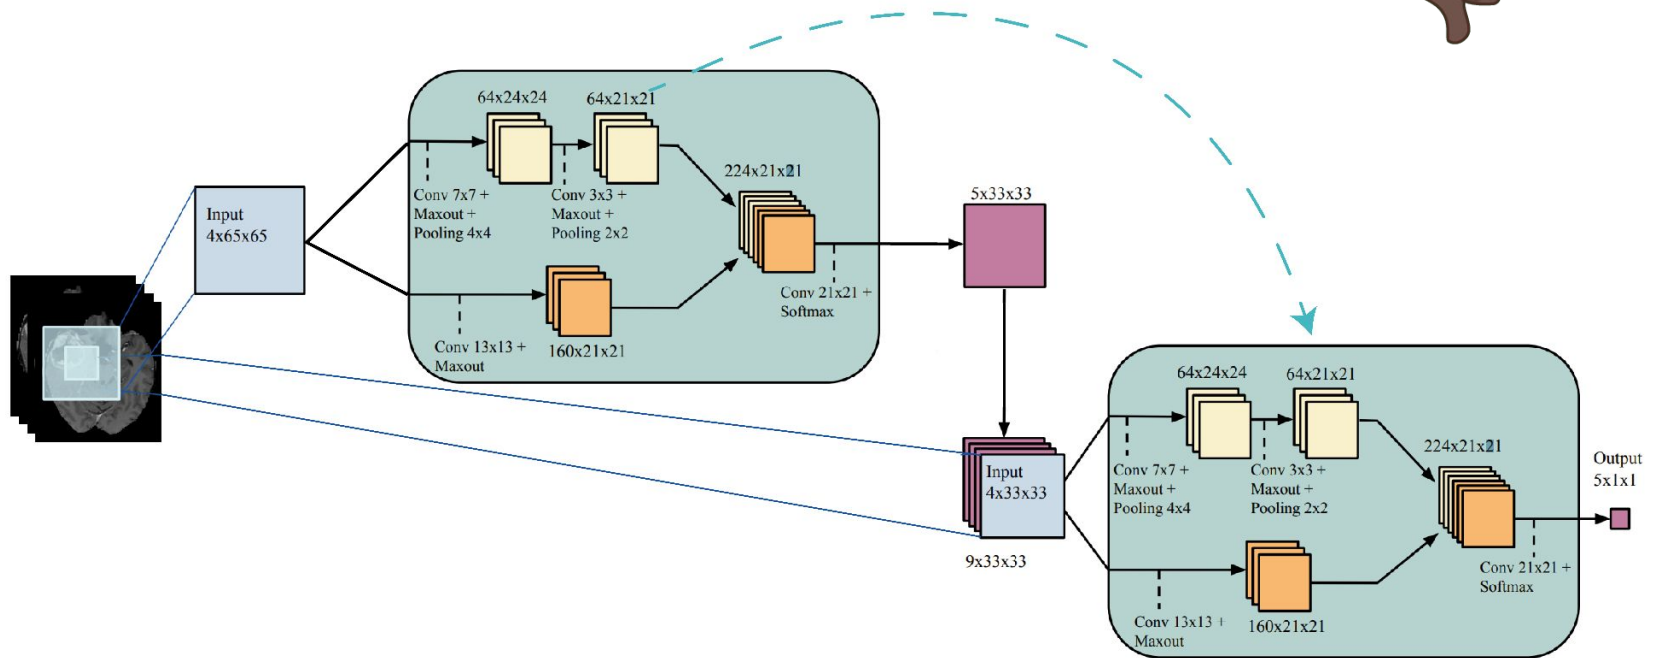


- 💡 Idea: are **all** learnt **parameters** actually **useful**?
- Test time → replace some layers with **identity** operation
- Same TL configs → 👎
- Keep local NN only → 👎
- Keep global NN only → 👎

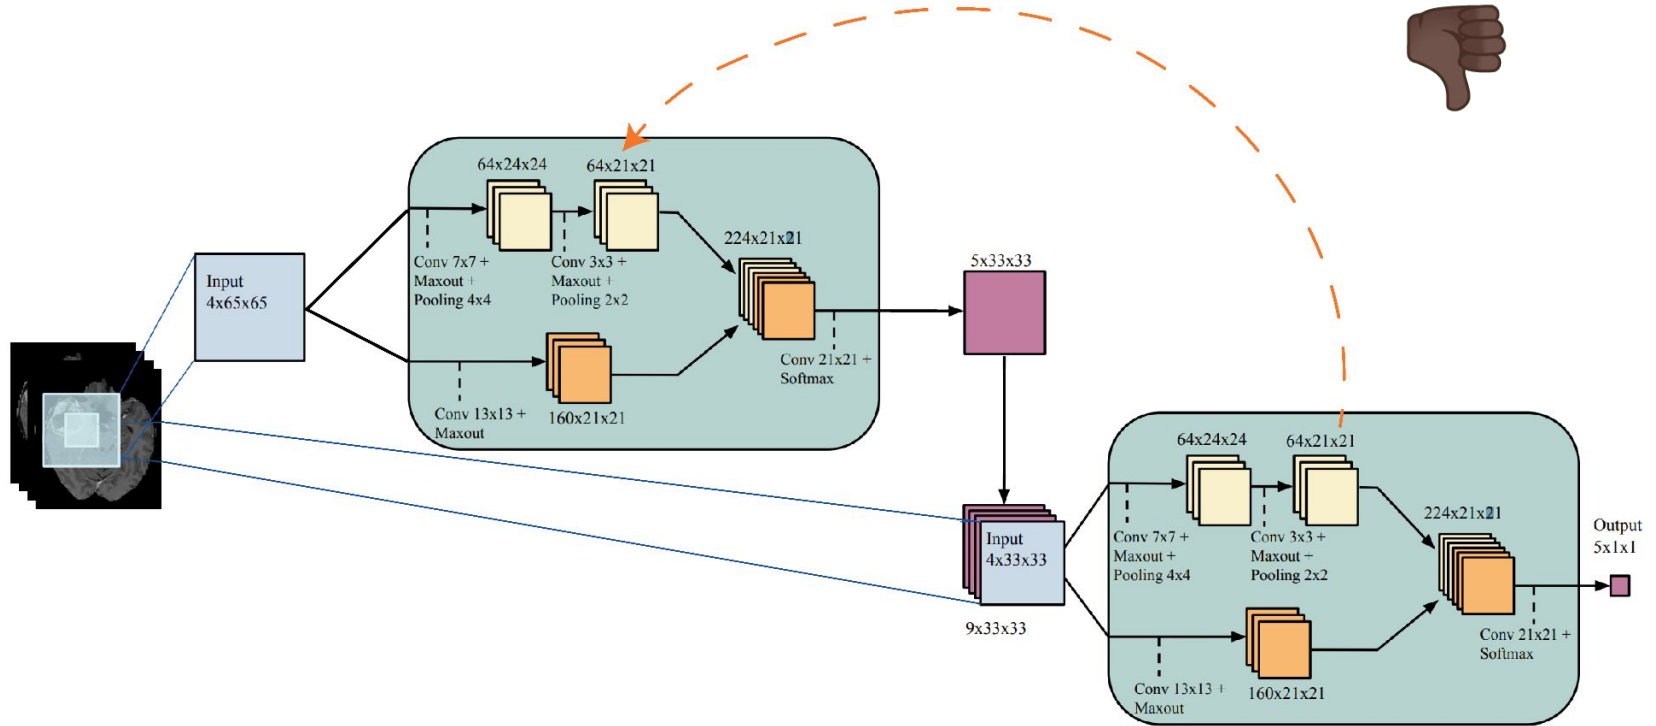
# Hypothesis: layer switch




# Hypothesis: layer copy (1/2)



# Hypothesis: layer copy (2/2)




# Bonus section



- Data **exploration** → NiBabel
- Data **cleaning**
- Data pre-processing → **parallelization**
- Data **normalization**
- Time and resource-**constrained** hyperparam tuning process
- Theory behind tested hypotheses
- **Weights and Biases** tracking 💜
- Custom-made train **dashboard** 💜

# Recap



- Brain and liver tumor segmentation via Transfer Learning
- Paper **novelties**:
  - TwoPathCNN, CNN concatenation, maxout activation, conv out layer, dual-stage training
- E2E brain tumor segmentation
- Multiple TL setups for liver tumor segmentation
  - **1/3** training **time** → **-1%** train and validation accuracy w.r.t. E2E
- **Challenges**:
  - Pre-processing, custom modules, hyperparam tuning, Transfer Learning
- Brain NN at test time **not invariant** w.r.t.
  - Parameter pruning, layer switch, layer copy